
Epidemic and Gossip Broadcast

Algorithmique répartie avancée - ARA
Master2

Luciana Arantes

22/09/2024

ARA: Broadcast - Partie 2

1

Epidemic Broadcast

Diffusion Epidémique

22/09/2024

ARA: Broadcast - Partie 2

2

Epidemic Broadcast

- **The broadcast algorithms that we have seen till now are not scalable**
 - They consider a set of processes known by all processes from the beginning.
- **Epidemic algorithms are effective solution for disseminating in large scale and dynamic systems.**
 - They do not provide deterministic broadcast guarantees but just make probabilistic claims about such guarantees.
- **An epidemic broadcast uses a randomized approach where all the participants in the protocol should collaborate in the same manner to disseminate information.**

22/09/2024

ARA: Broadcast - Partie 2

3

Epidemic Broadcast

- When a process p wishes to send a broadcast message, it selects k processes at random and sends the message to them
 - k is a typical configuration parameter called *fanout*.
- Upon receiving a message from p for the first time, a process q repeats the same procedure of p 's : q selects k gossip targets processes and forwards the message to them.
 - If a node receives the message twice, it simply discards the message
 - Each process needs to keep track of which messages it has already seen and delivered. The size of this buffer is also a scalable constraints.
- The step consisting of receiving a message and forwarding it is called a *round*.
 - An epidemic algorithm usually performs a *maximum number of rounds* r for each message.

22/09/2024

ARA: Broadcast - Partie 2

4

Epidemic broadcast

- **Epidemic broadcast can only be applied to applications that do not require full reliability.**
 - The cost of full reliability is usually not acceptable in large scale systems.
 - However, it is possible to build scalable randomized epidemic algorithms which provide good reliability guarantees.
 - It exhibit a very stable behavior even in the presence of failures.

Epidemic Broadcast

- **Parameters associated with the configuration of gossip protocols :**
 - **Fanout (k):** number of nodes that are selected as gossip targets by a node for each message that is received by the first time.
 - Tradeoff associated between desired reliability level and redundancy level of the protocol.
 - **Maximum rounds (r):** maximum number of times a given gossip message is retransmitted by nodes.
 - Each message carries a *round value*, which is increased each time the message is retransmitted.
 - **Modes :**
 - *Unlimited mode*: the parameter maximum round is undefined
 - *Limited mode*: the parameter maximum round is defined with a value greater than 0.
 - Higher value: higher reliability as well as message redundancy.

Epidemic Broadcast

- **Probabilistic Broadcast**
 - **Properties**
 - *Probabilistic validity*: There is a given probability such that for any two correct processes p_i and p_j , every message broadcast by p_i is eventually delivered by p_j with this probability.
 - *No duplication*: No message is delivered more than once by a process
 - *No creation*: If a message m is delivered by some process p_j , then m was previously broadcast by some process p_i .

Epidemic Broadcast

- **Strategies**
 - **Eager push approach**: Nodes send message to selected nodes as soon as they receive them for the first time
 - **Pull approach**: Periodically, nodes query random selected nodes for information about recently received messages. When they receive information about a message they did not received yet, they explicitly request the message to their neighbors.
 - **Lazy push approach**: When a node receives a message for the first time, it gossips only the message identifier. If a node receives a identifier of a message it has not received, it makes an explicitly pull request.
 - **Hybrid approach**: First phase uses a push gossip to disseminate a message in best-effort manner. A second phase of pull gossip is used to recover messages not received in the first phase.

Eager Push Epidemic Broadcast

Algorithm

```
Init :
    delivered = ∅

Epid_broadcast (m)
    gossip(self, m, maxrounds);


upon recv (pi, <src,m, r>)
    if (m ∉ delivered)
        delivered = delivered U {m}
        Epid_deliver(src,m)
    if (r > 0)
        gossip(self, m, maxrounds - 1);

Function chose-targets (ntargets)
    targets = ∅
    while ( | targets| < ntargets ) do
        candidate = random (Π)
        if ((candidate ∉ targets) and
            (candidate != self) )
            targets = targets U {candidate};
    return targets

procedure gossip (src,msg,round)
    for i ∈ chose-targets(fanout) do
        send (i, msg, round);
```

Eager Push Epidemic Broadcast

Execution example



Fanout = 3 ; Maxround = 3

Epidemic Broadcast

- **Ideally, one would like to have each participant to select gossip targets at random from the entire system, as shown in the previous example.**
 - Realistic if it is deployed within a moderate sized cluster.
 - Such approach is not scalable :
 - High memory cost to maintain full membership information.
 - High cost of ensuring the update of such information.
- **Solution:**
 - Gossip-based (epidemic) broadcast protocols rely on *partial view*, instead of full membership information.

Epidemic Broadcast : Partial view

- **Partial view**
 - A process just knows a small subset of the entire system membership, from which it can select nodes to whom relay gossip messages
 - The membership protocol establishes *neighboring* association among nodes.
 - It must maintain the partial view at each node in face of dynamic changes in the system membership.
 - Joining of new nodes, crashes of nodes, etc.
 - A partial view must be a tradeoff between *scalability* against *reliability*
 - Small views scale better, while large views reduce the probability that processes become isolated or that network partitions occur.
 - **Overlay**
 - Partial views of all nodes of the system define a graph

Epidemic Broadcast : Partial view

- **Partial View Properties : related to the graph properties of the overlay defined by the partial view of all nodes**
 - *Connectivity* : the overlay should be connected : there should be at least one path from each node to all other nodes.
 - *Degree Distribution* : number of edges of the node.
 - *In-degree* of node n : number of nodes that have n in their partial view. It provides a measure of *reachability*.
 - *Out-degree* of node n : number of nodes in n 's view: measure of the importance of that node to maintain the overlay.
 - *Average Path Length* : the average of all shortest paths between all pair of nodes in the overlay.


Epidemic Broadcast : Partial view

- **Strategies to maintain partial view**
 - *Reactive strategy* : a partial view only changes in response to some external event such as a joining of a node, a crash of a node, etc.
 - *Cyclic strategy* : A partial view is update every ΔT units of time, as a result of some periodic process that usually involves the exchange of information with one or more neighbors.
 - *Mixing strategy* : the partial view membership is included in the epidemic broadcast protocol
 - Whenever a process forwards a message, it also includes in it a set of processes it knows. Process that receives this message can update its own list of known processes.
 - It does not introduce extra communication to maintain membership.

Epidemic Broadcast : Partial view

CYCLON: Cyclic strategy

exchange of view periodically among neighbors (*shuffling operation*), at a fixed period ΔT .



Partial View 2 = {0,1,3,6,9}
Partial View 9 = {0,4,5,7}

Exchange :
2 -> 9 : {2,0,6}
9 -> 2 : {0,5,7}

Partial View 2 = {0,1,3,5,7}
Partial View 9 = {0,2,4,6}

Gossip protocol in ad hoc Networks

- **An ad hoc network is a multi-hop wireless network with no fixed infrastructure**
 - Node broadcasts a message which is received by all nodes within one hop (neighbors)
- **Gossiping protocol Gossip(p)[HHL06]**
 - A source node sends the message m with probability 1.
 - Upon reception of m
 - first time,
 - it broadcasts m with probability p
 - it discards m with probability $1-p$
 - Otherwise it discards m

Gossip protocol in ad hoc Networks

- **If the source has few neighbors, chance that none of them will gossip and the algorithm dies.**
 - Solution : Gossip (p, k)
 - Gossip with probability 1 for the k hops before continuing to gossip with probability p .
 - Gossip (1,1) is equivalent to flooding.
 - Gossip ($p, 0$) : even the source gossips with probability p .

Bibliographie

- R. Guerraoui, L. Rodrigues. *Reliable Distributed Programming*, Springer, 2006 .
- J. C. A. Leitão. *Gossip-based broadcast protocols*. Master thesis's. 2007.
- P. Eugster, R. Guerraoui, A. Kermarrec and L. Massoulié. *From Epidemics to Distributed Computing*. IEEE Computer, 37, pages 60-67.
- S. Voulgaris, D. Gavidia, M. Stten. *Cyclon : Inexpensive membership management for unstructured p2p overlays*. Journal of Network and System Management. Vol 13, pages 197-217, 2005.
- Z.J. Haas, J. Halpern, L. Li. *Gossip-Based Ad Hoc Routing*. IEEE Transactions on Network, Vol. 14, N. 13, pages 479-491, 2006