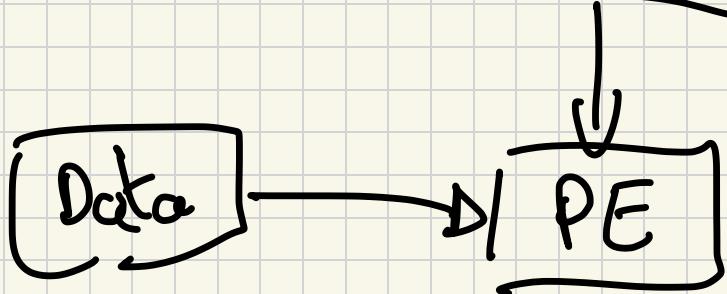


pierre @ jofiv. et

Flynn's

taxonomy
instructions

1966

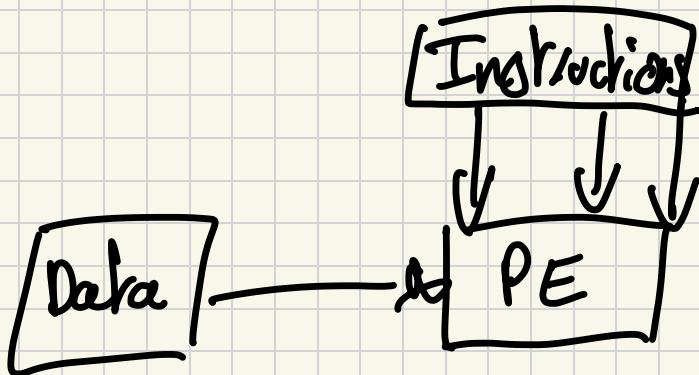


four kind of architectures :

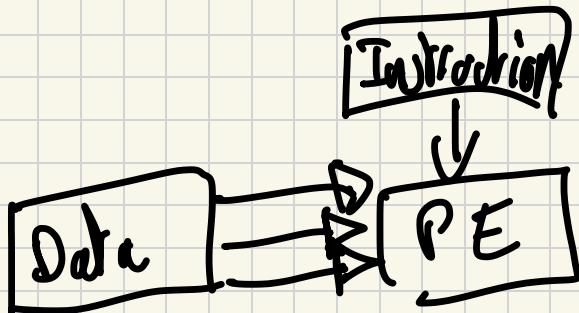
SISO : single instruction,
single data



MISO :



SIMD :



multiple instructions,
single data

single instruction,
multiple data
(Vectorization)

→ MIMD : multiple instructions,
multiple data

Very large and abstract set
of hardware, SPMOD :

single program multiple data
(Message Passing Interface)

Hardware considerations:

- * multiple PE, coordination:
synchronous or asynchronous
- * memory space:
shared or distributed
- * memory addressing:
local or global
(Remote Memory Addressing)

shared- memory

scaling	X
data integrity	X
performance optimisation	X
incremental parallelisation	✓
automatic	

distributed - memory

✓

✓

✓

X

☞ Matrix Storage

- graphs
- discrete systems

1) Dense matrices :

- * indexing (start at 0, or 1)
- * memory - layout (column-major or row-major)

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \Rightarrow \text{float } a[] = \left\{ 1, 2, 3, 4 \right\};$$

2) Sparse matrix (an operator
with a lot of zeros)

- coordinate format (COO)

n (rows), m (columns),

nnz (number of nonzero entries)
col (columns indices)

row (row indices), V (numerical values)

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 \\ 1 & 2 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 4 & 5 & 0 & 7 \end{bmatrix}$$

$n = 4$
 $m = 4$
 $nnz = 6$

$$col = \{0, 0, 1, 1, 2, 3\};$$

$$row = \{1, 3, 1, 3, 0, 3\};$$

$$V = \{1, 4, 2, 5, 1, 7\};$$

* How to compute $y = Ax$?

$$y_i = \sum_{j=1}^m A_{ij} \times j$$

In the dense case, need to perform $n \times m$ operations
But here we can exploit the sparsity of A

Single for loop :

```
for (int i = 0; i < n; ++i)  
    y[i] = 0;
```

```
for (int i = 0; i < nnz; ++i)
```

$$y[\text{row}[i]] += r[i] \cdot x[\text{col}[i]];$$

- compressed sparse row (CSR)

n (rows), m (columns)

row_disp (displacement of rows),
col (column indices), ✓ (numerical values)

row_displ is built recursively :

$\forall i \in [0; n-1] \text{ row_displ}[0] = 0;$
 $\text{row_displ}[i+1] = \text{row_displ}[i]$
+ number of non zero
of the $(i+1)^{\text{th}}$ row
assert ($\text{nnz} == \text{row_displ}[n]$);

$$A = \begin{bmatrix} 0 & 0 & 5 & 0 \\ 0 & 2 & 3 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$n = 4 \quad m = 4$$

$$nm = 4$$

$$col = \{ 2, 1, 2, 3 \} ;$$

$$row = \{ 5, 2, 3, 1 \} ;$$

$$row_displ = \{ 0, 1, 3, 4, 4 \} ;$$

How to compute $y = Ax$?

The number of operations is still $n \times n^2$

Use a nested loop

```
for(int i=0; i < n; ++i) {  
    y[i] = 0;  
    for(int j=row_dsp[i]; j < row_dsp[i+1];  
        ++j)  
        y[i] += v[j] * col_dsp[j];  
}
```

Again, this is difficult to parallelize (or even make efficient) because of cache misses and non-contiguous memory-access

- diagonal format (DIA)
 - n (rows), m (cols), v (numerical values)
 - $ndiag$ (number of diagonals),
 - $idiag$ (shifts w.r.t. the main diagonal)

$$A = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 7 & 2 & 0 & 6 \\ 0 & 8 & 3 & 0 \\ 10 & 0 & 9 & 4 \end{bmatrix}$$

(two-dimensional
array with
padding)

$$n = 4, \quad m = 4$$

$$\text{ndiag} = 4$$

$$\text{idiaj} = \{-3, -1, 0, 2\}$$

$$V = \begin{bmatrix} 0 & 0 & 0 & 10 \\ 0 & 7 & 8 & 9 \\ 1 & 2 & 3 & 4 \\ 0 & 6 & 0 & 0 \end{bmatrix}$$

* How to compute $y = Ax$?

for (int i=0; i<n; ++i)

$$y[i] = 0$$

for (int k=0; k<ndiag; ++k)

for (int j = max(0, -idiag(k));

j < min(n, n - idiag(k));

++j)

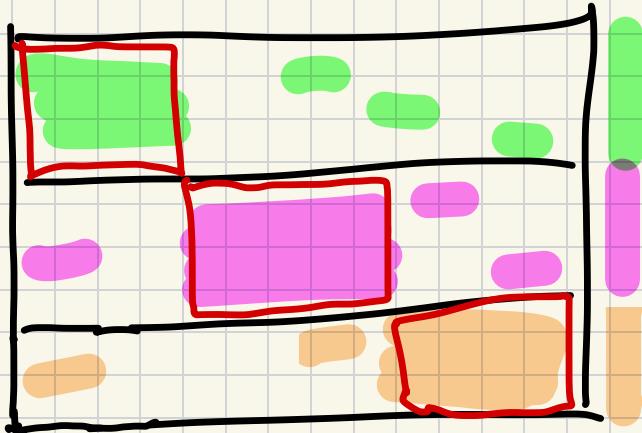
$$y[j] += v[k][j] \cdot x[j + idiag[k]];$$

Memory accesses in the inner-most loop are thus time contiguous. For vectorization, OpenMP should perform OK.

• What about distributed - memory parallelism?

We stick to the CSR format

$A =$



$y = Ax$

There is an issue w.r.t. the distribution of x

We could split according to columns but this doesn't fit the CSR format (\Rightarrow CSC)

If A is block-diagonal, we can distribute x ; i.e. we distribute A , and there is no additional cost.

Solving linear systems of equations

$$Ax = f$$

* Gaussian elimination

- For dense systems, complexity of order n^3
- For sparse systems, complexity can be lowered but it's most often Superlinear
 $(n^2 \text{ or more})$
- In sequential, it's straightforward to implement

There are some implementations for distributed - memory parallelism :

- MUMPS
- Pastix
- PARDISO
- SuperLU_DIST

* Iterative methods

- Much cheaper than direct methods
- But may not converge fast enough

We make the assumption that A^{-1} is too expensive to compute.

Instead we want to derive an iterative scheme:

$$x = x^{(k)} + e^{(k)}$$

exact solution ↓ error \rightarrow at iteration #
approximate solution ↑

What does $e^{(R)}$ satisfy?

$$\begin{aligned} Ax &= Ax^{(R)} + Ae^{(R)} \\ &= f \quad \Rightarrow \end{aligned}$$

$$Ae^{(R)} = f - Ax^{(R)}$$

We introduce $r^{(R)} = f - Ax^{(R)} \Rightarrow r^{(R)} = Ae^{(R)}$

In the end, we have $x = x^{(R)} + A^{-1}r^{(R)}$

Problem: We cannot compute $A^{-1}r^{(R)}$

Instead, introduce an approximation of A^{-1} , that we will denote M^{-1}

M^{-1} is what we call a preconditioner

There are various ways of building either M or M^{-1}

Basic preconditioners :

- * Jacobi : we pick M as the diagonal of A then, it is much cheaper to compute M^{-1}

In practice, the previous formula

$$x = x^{(R)} + A^{-1} r^{(R)}$$

becomes

the following iterative scheme :

$x^{(0)}$ is given

$$x^{(R+1)} = x^{(R)} + M^{-1} r^{(R)}$$

Alternatively, we can write:

$$\begin{aligned} x^{(R+1)} &= x^{(R)} + M^{-1} (f - Ax^{(R)}) \\ &= (I - M^{-1}A)x^{(R)} + M^{-1}f \end{aligned}$$

What about Jacobi preconditioning?

$$A = L + D + U$$

$$\begin{aligned}x^{(R+1)} &= (I - D^{-1}A)x^{(R)} + D^{-1}f \\&= \cancel{(I - D^{-1}(L+U) - I)} x^{(R)} + D^{-1}f \\&= -D^{-1}[(L+U)x^{(R)} - f]\end{aligned}$$

We need to update all the rows of

$$x^{(R+1)}$$

$$f_i \in [1; n] \quad x_i^{(k+1)} = \frac{-1}{A_{ii}} \left[\sum_{\substack{j=1 \\ j \neq i}}^n A_{ij} x_j^{(k)} - f_i \right]$$

$(R+1)$ on the LHS, (R) on the RHS, so we can update coefficient simultaneously

$$\text{if } A_{ii} = 0 \Rightarrow \frac{1}{A_{ii}} = 1$$

The convergence is not guaranteed but one can check $\|r^{(k)}\|$

At Gauss - Seidel method

$$M^{-1} = (L + D)^{-1}$$

$$x^{(k+1)} = (I - M^{-1}A)x^{(k)} + M^{-1}f$$

$$= (I - (L+D)^{-1}(L+D+U))x^{(k)} + (L+D)^{-1}f$$

$$= -(L+D)^{-1}[Ux^{(k)} - f]$$

$$(L+D)x^{(k+1)} = f - Ux^{(k)}$$

Let's pick a row number $i \in [1; n]$

$$\left((L+D) \times_j^{(R+1)} \right)_i = f_i - \left(U \times_j^{(R)} \right)_i$$

$$\sum_{j=1}^i A_{ij} x_j^{(R+1)} = f_i - \sum_{j=i+1}^n A_{ij} x_j^{(R)}$$

$$\sum_{j=1}^{i-1} A_{ij} x_j^{(R+1)} + A_{ii} x_i^{(R+1)} = f_i - \sum_{j=i+1}^n A_{ij} x_j^{(R)}$$

$$x_i^{(R+1)} = \frac{1}{A_{ii}} \left[f_i - \sum_{j=1+i}^n A_{ij} x_j^{(R)} - \sum_{j=1}^{i-1} A_{ij} x_j^{(R)} \right]$$

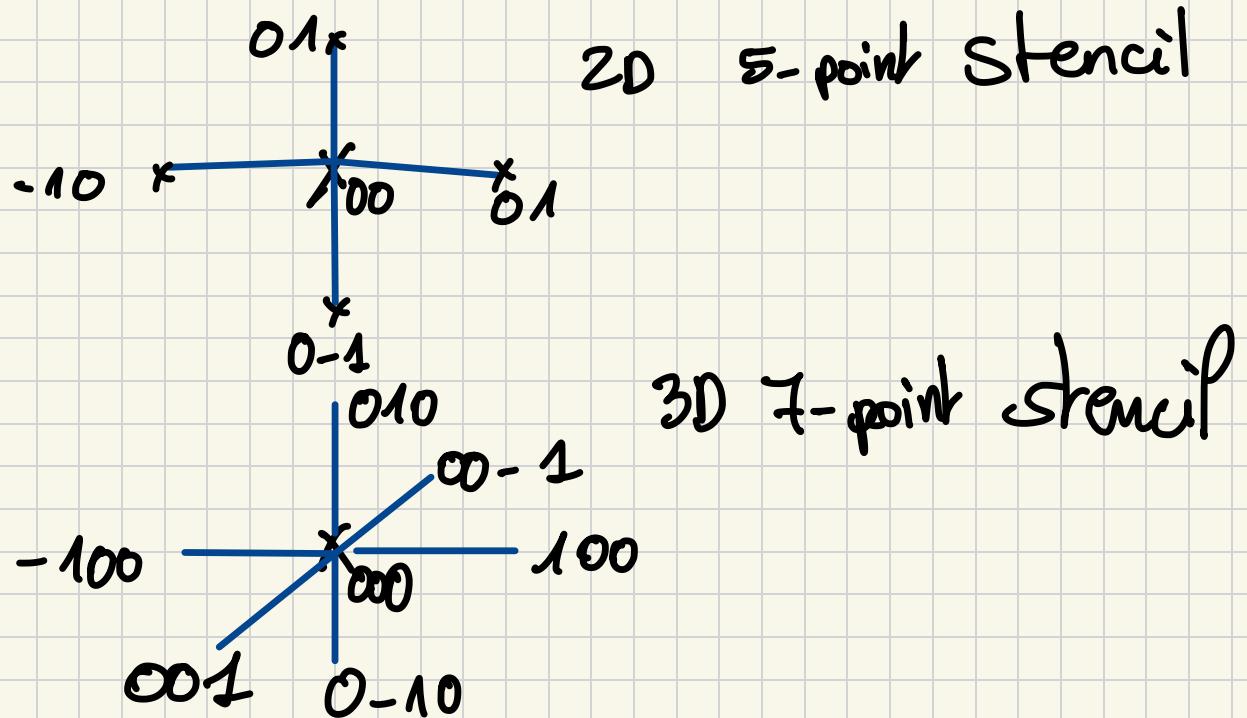
* Parallelizing the Gauss-Seidel method is much more difficult

With some assumptions, we can proceed as for the Jacobi method:

- * diagonal A
- * upper triangular A

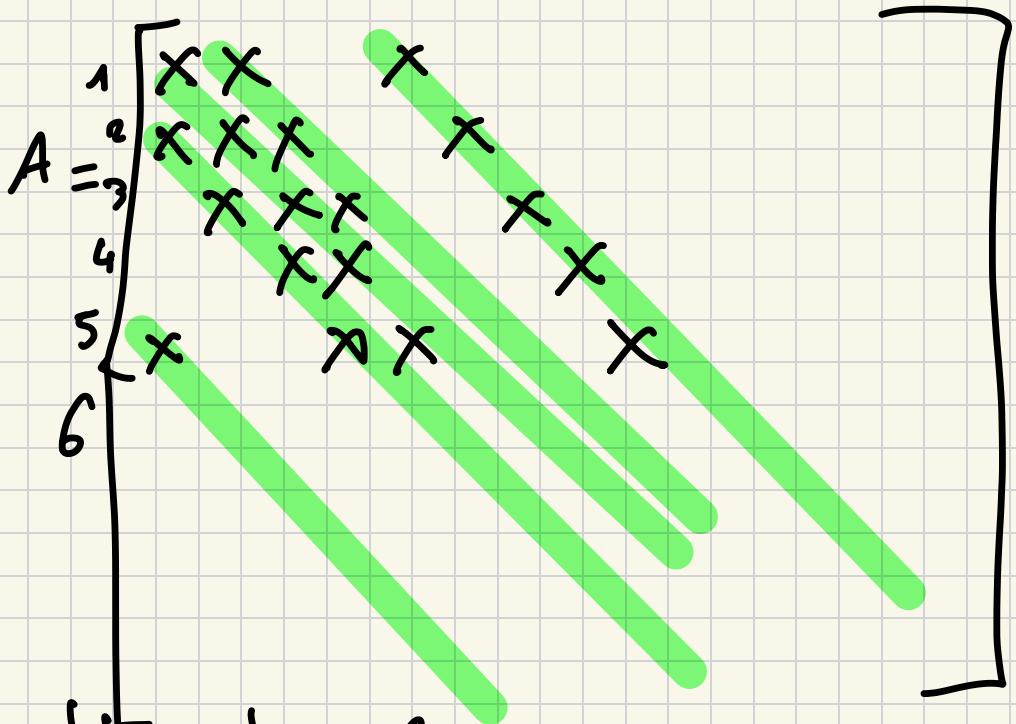
We can still make the preconditioner parallel using graph coloring techniques

Let us assume that A comes from
a stencil-based discretization:



	13	14	15	16
9		10	11	12
5				8
1	6	7		

Same number of
diagonals as the
number of point in the stencil

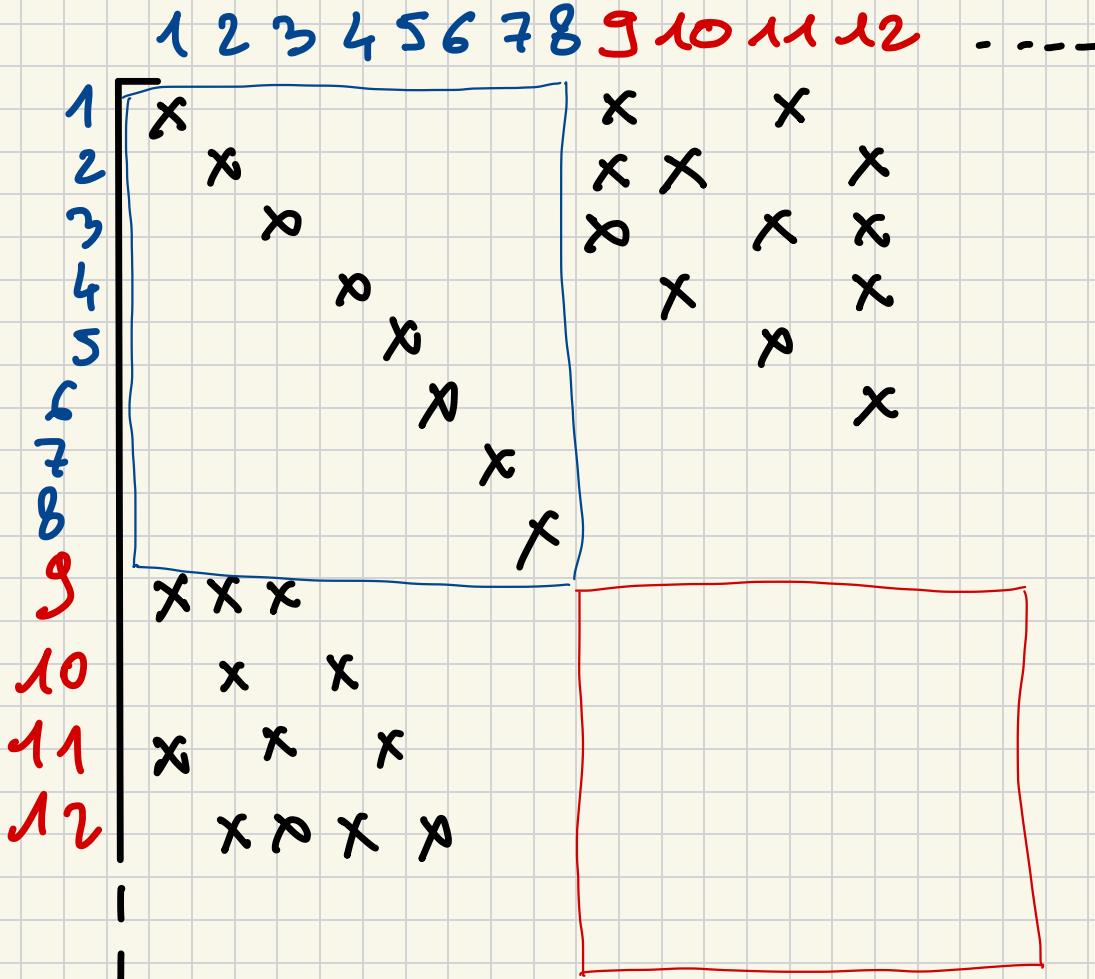


We need to renumber the unknowns
to exhibit possible parallelism:
red-black renumbering

⇒ we want to avoid having two grid
points in the same color (red or
black) interacting through the
stencil

15	7	16	8
5	13	6	14
11		12	4
1	9	2	10

thanks to the
RB numbering
diagonal blocks
are diagonal



The update formula for Gauss-Seidel is

$$x_i^{(R+1)} = \frac{1}{A_{ii}} \left[f_i - \sum_{j=1+i}^n A_{ij} x_j^{(R)} - \sum_{j=1}^{i-1} A_{ij} x_j^{(R+1)} \right]$$

* First we update all $x_i^{(R+1)}$ for
i in blue:

$$x_i^{(R+1)} = \frac{1}{A_{ii}} \left[f_i - \sum_{j=i+1}^n A_{ij} x_j^{(R)} \right]$$

This can be computed in parallel (just
as far Jacobi)

* Then, we can update the red unknowns

$$x_i^{(R+1)} = \frac{1}{A_{ii}} \left[f_i - \sum_{j=1}^{i-1} A_{ij} x_j^{(R+1)} \right]$$

Keep in mind that the blue unknowns were updated in the previous step so again, we can do this in parallel

- * There are more advanced
 - iterative methods
 - preconditioners

For iterative methods, we most often use Krylov methods, which are a specific kind of polynomial methods.

Remember the Cayley - Hamilton theorem

The characteristic polynomial of A

$$\chi(\lambda) = \det(\lambda I - A)$$

$$\chi(A) = 0$$

It is useful here because we can show that

$$\chi(X) = (-1)^{n+1} \det(A) + X q(X)$$

with $q \in K_{n \times n}[X]$

With the theorem, we get

$$\pi(A) = 0 = (-1)^{n+1} \det(A) I + A q(A)$$

$$A q(A) = (-1)^n \det(A) I$$

$$A \cdot q(A) = I \quad \det A \neq 0$$
$$\frac{(-1)^n}{(-1)^n \det(A)} q(A)$$

$$A^{-1} = \frac{q(A)}{(-1)^n \det(A)} \Rightarrow A^{-1} \text{ can be expressed as a polynomial of } A$$